

Evidence-Grounded Contract Clause Classification with Small Language Models

Anonymous ACL submission

Abstract

Commercial contract clauses exhibit substantial linguistic variation and semantic overlap, making fine-grained classification challenging for small language models (SLMs). We investigate whether retrieving semantically similar labeled clauses can provide task-specific evidence at inference time and improve classification without increasing model capacity. Using the Contract Understanding Atticus Dataset (CUAD), we formulate clause classification as a retrieval-conditioned prediction task in which relevant training clauses are retrieved and provided as in-context demonstrations to a locally deployed SLM. We evaluate the same classifier with and without retrieval under a contract-disjoint, leakage-controlled protocol, thereby isolating the contribution of retrieved evidence. Retrieval conditioning substantially improves performance, increasing macro-F1 from 31.55% to 56.51% on held-out clauses. Analysis of retrieval quality further shows that relevant clause categories are frequently represented among the retrieved examples. These findings suggest that retrieval can function as instance-level task grounding, enabling SLMs to leverage external contractual evidence rather than relying solely on parametric knowledge. The results support retrieval-conditioned SLMs as a promising direction for evidence-grounded and locally deployable legal NLP.

Keywords: Legal NLP, Contract Review, CUAD, Retrieval-Augmented Generation, Clause Classification, Evidence Grounding, Vector Retrieval, Qdrant, Small Language Models

1 Introduction

Commercial contracts encode rights, obligations, restrictions, and risk allocations through language that is both domain-specific and highly variable across agreements. Automatically identifying such provisions is therefore a central problem in legal

NLP, with applications to contract review, due diligence, compliance analysis, and legal information extraction. Benchmarks such as the Contract Understanding Atticus Dataset (CUAD) (Hendrycks et al., 2021) have enabled systematic evaluation of models for identifying legally salient contract clauses, while subsequent work has explored contextual modeling and large language models (LLMs) for contract understanding.

Recent work has explored the use of large language models for contract analysis, including structured question answering over specific contractual clauses (Roegiest et al., 2023). However, legal documents pose particular challenges for language models due to specialized terminology, complex syntax, and long-context requirements (Hua et al., 2023; Blair-Stanek et al., 2024). Studies evaluating general-purpose LLMs on legal contract classification have also reported limitations compared with models adapted to the legal domain (Jayakumar et al., 2023). LLMs have also been applied to practical contract-analysis tasks such as comparing contracts against templates and identifying deviations between contractual provisions (Narendra et al., 2024). These limitations motivate investigating whether external contractual evidence retrieved from a legal knowledge base can improve the performance of a small language model (SLM) for clause classification.

Despite this progress, clause classification remains challenging for small, locally deployable SLMs. A target provision may be difficult to classify from its text and a set of abstract category names alone, particularly when clause categories are semantically related or expressed using heterogeneous contractual language. Larger models may encode more task-relevant knowledge in their parameters, but relying on increasingly large models is not always desirable in legal settings where computational cost, data governance, and local deployment are important. This motivates a

different question: **can relevant contractual examples retrieved at inference time provide the task-specific evidence that an SLM lacks in its parametric representation?**

Retrieval offers a natural mechanism for addressing this problem. Rather than asking a classifier to infer a clause category in isolation, a retrieval-conditioned model can compare the target provision with semantically related, previously labeled contractual clauses. In this setting, retrieval is not merely a mechanism for locating documents. The retrieved examples provide instance-level task grounding by exposing the model to concrete realizations of legal categories at inference time. This perspective is particularly relevant to contract analysis, where similar legal functions can be expressed through substantially different lexical and syntactic forms.

Prior work has established the importance of evidence and contextual information for contract understanding. ContractNLI (Koreeda and Manning, 2021) explicitly associates contractual inference with supporting evidence, while ConReader (Xu et al., 2022) models relations extending beyond isolated clauses. Retrieval-augmented approaches have further demonstrated the value of grounding legal question answering in relevant legal provisions (Redelaar et al., 2024), while recent work has highlighted retrieval reliability as an important source of downstream error (Reuter et al., 2025). However, less is known about the extent to which retrieved labeled contractual examples can improve fine-grained clause classification for capacity-constrained SLMs under a controlled inference setting. In particular, separating the contribution of retrieved evidence from changes in the underlying language model is important for determining whether improvements arise from retrieval itself.

We study this question through retrieval-conditioned clause classification on CUAD. We construct the retrieval corpus exclusively from training contracts and retrieve semantically similar labeled clauses for each unseen target clause. These examples are supplied as in-context demonstrations to a locally deployed Qwen2.5-1.5B classifier. To isolate the effect of retrieval, we compare the same SLM with and without retrieved examples while holding the evaluation instances, label space, and generation settings constant. Contracts are partitioned before clause-level restructuring, ensuring that retrieval cannot draw evidence from

the contract containing an evaluation instance, and annotation labels are excluded from retrieval representations. The key contributions of this work are as follows.

- We formulate fine-grained contract clause classification as evidence-conditioned inference, where semantically similar labeled clauses provide task-specific grounding for an SLM at inference time.
- We demonstrate, under a contract-disjoint and leakage-controlled setting, that retrieved contractual evidence substantially improves the same SLM, increasing macro-F1 from 31.55% to 56.51%.

The rest of the paper is organized as follows. Section 2 reviews related work. Section 3 presents the retrieval-conditioned methodology. Section 4 describes the experimental setup. Section 5 reports the results, and Section 6 concludes the paper with limitations and future work.

2 Related Work

Prior work related to this study falls into three connected lines: evidence-centered contract understanding, retrieval-grounded legal NLP, and LLM-based clause analysis. Table 1, summarizes the studies that directly inform our experimental setting.

These studies establish two properties that are central to our setting. First, contractual decisions often depend on evidence or context beyond the target clause itself (Hendrycks et al., 2021; Koreeda and Manning, 2021; Xu et al., 2022). Second, retrieval can provide such context but must itself be treated as an uncertain component of the pipeline: legal RAG has benefited from attributed retrieval (Redelaar et al., 2024), while retrieval failures can propagate to downstream predictions (Reuter et al., 2025). ContractEval further shows that LLM performance varies substantially across models on clause-level legal analysis (Liu et al., 2025).

Taken together, these trends expose a specific gap: it remains unclear whether improvements in fine-grained contract classification can be obtained by supplying a capacity-constrained language model with relevant labeled contractual evidence at inference time, rather than increasing model capacity. Addressing this gap requires an experimental design that (i) retrieves evidence

Table 1: Prior work on evidence, contractual context, retrieval, and language-model-based contract analysis.

Study	Task	Relevant contribution
CUAD (Hendrycks et al., 2021)	Contract review	Expert-annotated clause-level contract review across 41 categories, providing a benchmark for identifying legally salient provisions.
ContractNLI (Koreeda and Manning, 2021)	Contract inference	Document-level inference paired with explicit identification of supporting contractual evidence.
ConReader (Xu et al., 2022)	Clause extraction	Models long-range context, term-definition relations, and related contractual provisions beyond the target clause.
Redelaar et al. (Redelaar et al., 2024)	Legal QA	Retrieval-augmented legal QA with attribution to supporting legal provisions.
Reuter et al. (Reuter et al., 2025)	Legal RAG	Studies retrieval reliability and identifies document-level retrieval mismatch as a major failure mode.
ContractEval (Liu et al., 2025)	LLM clause analysis	Benchmarks proprietary and open-source LLMs for clause-level legal-risk identification on CUAD.

exclusively from training contracts, (ii) prevents document- and label-level leakage, (iii) evaluates retrieval quality independently from downstream prediction, and (iv) compares evidence-free and evidence-conditioned inference while holding the underlying classifier and evaluation setting fixed. Accordingly, we formulate clause classification as a retrieval-conditioned prediction problem, in which semantically related labeled clauses are retrieved from a contract-disjoint training corpus and supplied as in-context demonstrations to the same small language model. The following section describes this methodology.

3 Methodology

The proposed framework uses semantic retrieval to provide relevant contractual evidence before clause classification. Given an input clause, the retrieval component identifies similar clauses from the training corpus, and Qwen2.5- 1.5B uses these labeled examples to assign the final CUAD category. The separation between retrieval and prediction also permits the contribution of retrieved evidence to be evaluated independently.

3.1 Task Formulation

Let q denote a target contractual clause and $\mathcal{Y} = \{y_1, \dots, y_m\}$ the set of candidate clause categories. Standard clause classification predicts a label using only the target clause and the predefined label space:

$$\hat{y}_{\text{base}} = g_{\phi}(q, \mathcal{Y}), \quad (1)$$

where g_{ϕ} denotes the language-model classifier.

We instead formulate clause classification as *retrieval-conditioned inference*. Let $\mathcal{D}_{\text{train}} =$

$\{(c_i, y_i)\}_{i=1}^N$ denote the training-only corpus of labeled contractual clauses, and let $f_{\theta}(\cdot)$ be a dense clause encoder. For a target clause q , the retrieval stage constructs an evidence set

$$E_q = \text{TopK}_{(c_i, y_i) \in \mathcal{D}_{\text{train}}} \text{sim}(f_{\theta}(q), f_{\theta}(c_i)), \quad (2)$$

where $\text{sim}(\cdot, \cdot)$ denotes cosine similarity and E_q contains the k highest-ranked labeled clauses.

The retrieved evidence is subsequently provided to the same classifier as in-context demonstrations:

$$\hat{y}_{\text{ret}} = g_{\phi}(q, E_q, \mathcal{Y}). \quad (3)$$

Thus, the baseline and retrieval-conditioned settings differ only in the availability of E_q ; the underlying classifier remains unchanged. This formulation allows us to test whether retrieved contractual examples provide instance-level task grounding beyond the parametric knowledge of a small language model.

Figure 1, illustrates the resulting framework, consisting of dense retrieval over a training-only clause index, construction of an evidence set from the top- k labeled clauses, and evidence-conditioned clause prediction.

3.2 Evidence Retrieval

The retrieval corpus $\mathcal{D}_{\text{train}}$ contains labeled clauses derived exclusively from the training contracts. Each clause c_i is mapped to a dense representation $f_{\theta}(c_i)$ and stored in a vector index. Importantly, clause labels and annotation answers are excluded from the embedding input, ensuring that retrieval is based solely on contractual text. For a target clause q , we compute $f_{\theta}(q)$ and rank training clauses by




Figure 1: Overview of the retrieval-conditioned clause classification framework. A target clause is encoded and used to retrieve the top- k semantically similar labeled clauses from a training-only index. The retrieved clauses serve as in-context evidence for a small language model, which predicts the corresponding CUAD clause category.

cosine similarity:

$$s(q, c_i) = \frac{f_\theta(q)^\top f_\theta(c_i)}{\|f_\theta(q)\| \|f_\theta(c_i)\|}. \quad (4)$$

The top $k = 5$ clauses and their associated labels form the evidence set E_q used for downstream classification. Because the index contains only training contracts, no clause from a held-out contract can be retrieved as evidence.

3.3 Evidence-Conditioned Classification

The retrieved evidence set $E_q = \{(c_j, y_j)\}_{j=1}^k$ is incorporated into the classifier prompt as a set of labeled in-context demonstrations. Each retrieved pair provides a concrete example of contractual language associated with a candidate clause category. Given an instruction I , the permitted label set $\mathcal{Y}$, and target clause q , the evidence-conditioned prompt is defined as

$$P(q, E_q, \mathcal{Y}) = [I; E_q; \mathcal{Y}; q]. \quad (5)$$

The language model then predicts

$$\hat{y} = g_\phi(P(q, E_q, \mathcal{Y})), \quad (6)$$

where the output is restricted to a category in $\mathcal{Y}$. In contrast, the evidence-free baseline receives the same instruction, label space, and target clause but excludes E_q . This design enables the effect

of retrieved contractual evidence to be assessed without modifying the parameters of the underlying classifier.

3.4 Controlled Evaluation and Leakage Prevention

To isolate the effect of retrieved evidence, we compare two inference conditions using the same classifier:

$$\hat{y}_{\text{base}} = g_\phi(q, \mathcal{Y}), \quad (7)$$

$$\hat{y}_{\text{ret}} = g_\phi(q, E_q, \mathcal{Y}). \quad (8)$$

The model, evaluation instances, label space, and decoding configuration are held fixed; the only experimental difference is the inclusion of the retrieved evidence set E_q .

To prevent information leakage, contracts are partitioned before clause-level restructuring, and the retrieval index is constructed exclusively from training contracts. Consequently, clauses from held-out contracts cannot appear as retrieved evidence. Annotation labels are also excluded from the embedding representations, ensuring that retrieval similarity is determined only by clause text.

This controlled design separates the contribution of retrieval from changes to the underlying classifier and provides a direct test of whether retrieved contractual evidence improves clause classification.

4 Experimental Setup

In this section, we describe the dataset and experimental procedure in detail.

4.1 Data

We use CUAD v1, an expert-annotated benchmark for commercial contract review containing 510 contracts and 41 annotation categories (Hendrycks et al., 2021). These 41 categories comprise five contract-level metadata fields and 36 substantive clause categories. The five metadata fields are excluded from the classification label space, so the experimental task considers 36 clause categories.

To prevent document-level leakage, the dataset is partitioned at the contract level before clause-level restructuring. Using random seed 42, 410 contracts are assigned to training and 100 to the held-out set, with zero overlap between the two partitions. We treat each annotated clause as one classification instance, with its CUAD category assigned as the class label. The training contracts are then converted into clause-level records containing the source contract, clause category, clause text, annotation answer, and metadata flag. Annotation answers remain separate from clause text and are not used as embedding input. The resulting training representation contains 7,004 clause records, which form the retrieval corpus.

4.2 Experimental Configuration

All classification configurations are evaluated on the same 200 held-out clauses sampled with seed 42. The evaluation subset contains instances from 31 of the 36 clause categories; five clause categories have zero support in this 200-clause sample. The retrieval corpus contains clauses from the 410 training contracts only. Each test clause is encoded with embed-english-v3.0, and the five most similar training clauses are retrieved from a local Qdrant collection using cosine similarity.

Both classification conditions use the same Qwen2.5-1.5B model served through Ollama with deterministic generation (temperature=0). The Qwen-only baseline receives the target clause and the CUAD label set. The retrieval-augmented condition additionally receives the five retrieved clauses and their labels as few-shot examples. Thus, the model, evaluation instances, label space, and generation settings are held constant across the comparison.

4.3 Evaluation Protocol

Retrieval quality is measured using Top-1 accuracy, Recall@3, Recall@5, and mean reciprocal rank (MRR). Classification performance is evaluated using accuracy and macro-F1; macro-precision and macro-recall are additionally reported for the baseline analysis. The same 200 held-out clauses are used throughout, allowing the effect of retrieved contractual context to be assessed under a matched evaluation protocol. Because five of the 36 clause categories have zero support in this evaluation subset, per-category recall is reported for the 31 categories represented in the sample; the five zero-support categories are retained in the category inventory but do not contribute meaningful per-category recall estimates.

5 Results

The outcome of our evaluation is described as follows.

5.1 Clause Classification

Table 2 shows the matched comparison between Qwen2.5-1.5B without retrieved evidence and the same model conditioned on five retrieved contractual examples.

The Qwen-only baseline achieves 46.50% accuracy and 31.55% macro-F1. With the same model, test clauses, label space, and generation settings, adding five retrieved contractual examples increases the reported accuracy to 70.35% and macro-F1 to 56.51%. The corresponding absolute gains are 23.85 and 24.96 percentage points. Because the underlying classifier is unchanged, this comparison directly evaluates the effect of supplying retrieved contractual context as additional input. The retrieval-only Top-1 baseline achieves 65.00% accuracy, while evidence-conditioned classification reaches 70.35%, indicating that the language model provides additional predictive value beyond directly adopting the label of the nearest retrieved clause.

Table 2: Clause-classification performance on the same 200 held-out clauses.

Configuration	Accuracy	Macro-F1
Baseline model	46.50	31.55
Retrieval-only (Top-1)	65.00	50.31
Proposed	70.35	56.51
Absolute gain	+23.85	+24.96

5.2 Retrieval Performance

We separately evaluate the quality of the retrieval stage, independent of the downstream classifier. The retrieval metrics include Top-1 accuracy, Recall@3, Recall@5, and mean reciprocal rank (MRR) for the retrieval component on the same 200 held-out clauses. Because the baseline configuration performs no retrieval, these metrics apply only to the retrieval stage and do not have a corresponding baseline-model value; they are reported to characterize the quality of the evidence supplied to the proposed model, rather than to compare against a baseline.

The retrieval stage achieves 65.00% Top-1 accuracy, 82.00% Recall@3, 87.00% Recall@5, and an MRR of 0.7421 on the same evaluation sample. The 22-point gap between Top-1 accuracy and Recall@5 indicates that relevant category evidence is often retrieved even when it is not ranked first. This supports the use of multiple retrieved demonstrations rather than nearest-neighbor prediction alone. As reported in Section 5.1, the retrieval-only Top-1 label alone reaches 65.00% classification accuracy, while the full evidence-conditioned model reaches 70.35% indicating that the language model itself contributes additional predictive value beyond simply adopting the label of the nearest retrieved clause, and that retrieval therefore provides a broader candidate pool from which the downstream classifier can select.

5.3 Error Analysis

The aggregate metrics in Table 2 show the overall effect of retrieved evidence, but not which clause categories drive the remaining errors. Figure 2 shows the confusion matrices for the two classification configurations, and Table 3 reports per-category recall for the categories represented in the evaluation sample, with support counts shown explicitly. Figure 3 visualizes the six categories most affected by retrieval.


Retrieval yields the largest recall gains for Cap On Liability (14.3% $\rightarrow$ 100%, +85.7pp) and Ip Ownership Assignment (0% $\rightarrow$ 66.7%, +66.7pp), followed by Non-Transferable License (14.3% $\rightarrow$ 71.4%, +57.1pp), Revenue/Profit Sharing (36.4% $\rightarrow$ 90.9%, +54.5pp), Change Of Control (16.7% $\rightarrow$ 66.7%, +50.0pp), and Minimum Commitment (27.3% $\rightarrow$ 72.7%, +45.5pp). These categories share a common property: their defining language is often expressed through varied phrasing across

contracts, so a small model benefits substantially from seeing concrete labeled examples at inference time.


Some categories remain difficult even with retrieved evidence. Post-Termination Services improves from 0% to 44.4% recall but remains one of the weaker categories, and Rofr/Rofo/Rofn improves only marginally (0% $\rightarrow$ 12.5%), suggesting that retrieval provides limited disambiguating signal for these categories. In contrast, License Grant recall is unchanged at 73.3% under both configurations, though its precision improves substantially (35.5% $\rightarrow$ 57.9%) and its F1 score rises from 47.8% to 64.7%; qualitative inspection shows that under retrieval, fewer non-License-Grant clauses are misclassified as License Grant, even though the recall of true License Grant clauses itself does not change.

A small number of categories show a recall decrease under retrieval: Anti-Assignment falls from 100% to 75%, and Warranty Duration falls from 66.7% to 33.3%. Both categories have limited support (12 and 3 clauses respectively), so these declines should be interpreted cautiously rather than as evidence of a systematic failure mode.

The paired matrices show how retrieved context changes the distribution of errors across the CUAD categories. The comparison is used to identify categories that remain difficult despite additional contractual evidence and to distinguish persistent class ambiguity from errors that are reduced by retrieval.



(a) Baseline: Qwen-only



(b) Proposed: Qwen with retrieved evidence

Figure 2: Confusion matrices for clause classification on the same 200 held-out clauses. (a) Baseline model without retrieved evidence. (b) Proposed model with five retrieved labeled contractual examples.

Table 3: Per-category recall on the 200 held-out clauses for the baseline (Qwen only) and proposed (Qwen with retrieved evidence) configurations, across the 36 CUAD clause categories.

Clause Category	Support	Baseline Recall (%)	Proposed Recall (%)	Δ (pp)
Affiliate License-Licensee	2	50.0	0.0	-50.0
Anti-Assignment	12	100.0	75.0	-25.0
Audit Rights	18	66.7	100.0	+33.3
Cap On Liability	7	14.3	100.0	+85.7
Change Of Control	6	16.7	66.7	+50.0
Competitive Restriction Exception	3	0.0	0.0	+0.0
Covenant Not To Sue	1	0.0	100.0	+100.0
Exclusivity	12	33.3	66.7	+33.3
Governing Law	10	100.0	100.0	+0.0
Insurance	18	77.8	88.9	+11.1
Ip Ownership Assignment	6	0.0	66.7	+66.7
Irrevocable Or Perpetual License	3	33.3	33.3	+0.0
Joint Ip Ownership	2	100.0	100.0	+0.0
License Grant	15	73.3	73.3	+0.0
Liquidated Damages	2	50.0	50.0	+0.0
Minimum Commitment	11	27.3	72.7	+45.5
Most Favored Nation	1	0.0	0.0	+0.0
No-Solicit Of Customers	3	66.7	66.7	+0.0
No-Solicit Of Employees	0	—	—	—
Non-Compete	3	0.0	66.7	+66.7
Non-Transferable License	7	14.3	71.4	+57.1
Notice Period To Terminate Renewal	5	0.0	0.0	+0.0
Post-Termination Services	9	0.0	44.4	+44.4
Price Restrictions	0	—	—	—
Renewal Term	7	100.0	100.0	+0.0
Revenue/Profit Sharing	11	36.4	90.9	+54.5
Rofr/Rofn/Rofn	8	0.0	12.5	+12.5
Source Code Escrow	4	75.0	100.0	+25.0
Termination For Convenience	6	16.7	50.0	+33.3
Third Party Beneficiary	1	0.0	100.0	+100.0
Uncapped Liability	3	0.0	0.0	+0.0
Volume Restriction	1	0.0	100.0	+100.0
Warranty Duration	3	66.7	33.3	-33.3
Affiliate License-Licensor	0	—	—	—
Non-Disparagement	0	—	—	—
Unlimited/All-You-Can-Eat-License	0	—	—	—




Figure 3: Per-category recall comparison between the baseline Qwen-only model and the proposed retrieval-conditioned model for the six clause categories most affected by retrieved contractual evidence.

6 Conclusion

This work investigated whether retrieved contractual evidence can improve fine-grained clause classification for small language models without increasing model capacity. Under a contract disjoint and leakage-controlled evaluation, conditioning the same classifier on semantically similar labeled clauses increased macro-F1 from 31.55% to 56.51%, demonstrating a substantial benefit from retrieval while holding the underlying model and inference configuration fixed. These findings provide

initial evidence that retrieval as a form of *instance-level task grounding*, enabling a small language model to exploit task-specific evidence rather than relying solely on parametric knowledge. This capability is particularly relevant to contract-analysis settings where confidential or commercially sensitive documents may motivate local processing and reduced dependence on externally hosted language models. Although the present study is limited to a single retrieval configuration and a subset of held-out CUAD clauses, the results motivate further investigation of retrieval-conditioned SLMs as an efficient and locally deployable approach to legal NLP, including evaluation of alternative retrieval strategies, broader contract-analysis tasks, and explicit privacy and deployment requirements.

Limitations

This study evaluates the proposed approach on a sample of 200 held-out CUAD clauses. While the contract-disjoint evaluation provides a controlled setting for measuring the effect of retrieved evidence, evaluation on the complete held-out set would provide stronger estimates across clause categories. The experiments consider a fixed retrieval configuration using dense semantic retrieval with cosine similarity and $K = 5$. Accordingly, the reported results capture the effectiveness of the retrieval-conditioned framework under this setting; exploring alternative retrieval strategies and values of K remains an avenue for future work. Finally, the current evaluation focuses on clause classification within CUAD. The extent to which retrieval-conditioned small language models generalize to other contract-analysis tasks and legal datasets remains to be investigated.

References

- Andrew Blair-Stanek, Nils Holzenberger, and Benjamin Van Durme. 2024. [BLT: Can large language models handle basic legal text?](#) In *Proceedings of the Natural Legal Language Processing Workshop 2024*, pages 216–232, Miami, FL, USA. Association for Computational Linguistics.
- Dan Hendrycks, Collin Burns, Anya Chen, and Spencer Ball. 2021. [CUAD: An expert-annotated NLP dataset for legal contract review](#). In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*.
- Wenyue Hua, Yuchen Zhang, Zhe Chen, Josie Li, and Melanie Weber. 2023. [Mixed-domain language modeling for processing long legal documents](#). In *Proceedings of the Natural Legal Language Processing Workshop 2023*, pages 51–61, Singapore. Association for Computational Linguistics.
- Thanmay Jayakumar, Fauzan Farooqui, and Luqman Farooqui. 2023. [Large language models are legal but they are not: Making the case for a powerful Legal-LLM](#). In *Proceedings of the Natural Legal Language Processing Workshop 2023*, pages 223–229, Singapore. Association for Computational Linguistics.
- Yuta Koreeda and Christopher Manning. 2021. [ContractNLI: A dataset for document-level natural language inference for contracts](#). In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 1907–1919, Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Shuang Liu, Zelong Li, Ruoyun Ma, Haiyan Zhao, and Mengnan Du. 2025. [ContractEval: Benchmarking LLMs for clause-level legal risk identification in commercial contracts](#). In *Proceedings of the Natural Legal Language Processing Workshop 2025*, pages 291–291, Suzhou, China. Association for Computational Linguistics.
- Savinay Narendra, Kaushal Shetty, and Adwait Ratnaparkhi. 2024. [Enhancing contract negotiations with LLM-based legal document comparison](#). In *Proceedings of the Natural Legal Language Processing Workshop 2024*, pages 143–153, Miami, FL, USA. Association for Computational Linguistics.
- Felicia Redelaar, Romy Van Drie, Suzan Verberne, and Maaïke De Boer. 2024. [Attributed question answering for preconditions in the Dutch law](#). In *Proceedings of the Natural Legal Language Processing Workshop 2024*, pages 154–165, Miami, FL, USA. Association for Computational Linguistics.
- Markus Reuter, Tobias Lingenberg, Ruta Liepina, Francesca Lagioia, Marco Lippi, Giovanni Sartor, Andrea Passerini, and Burcu Sayin. 2025. [Towards reliable retrieval in RAG systems for large legal datasets](#). In *Proceedings of the Natural Legal Language Processing Workshop 2025*, pages 17–30, Suzhou, China. Association for Computational Linguistics.
- Adam Roegiest, Radha Chitta, Jonathan Donnelly, Maya Lash, Alexandra Vtyurina, and Francois Longtin. 2023. [Questions about contracts: Prompt templates for structured answer generation](#). In *Proceedings of the Natural Legal Language Processing Workshop 2023*, pages 62–72, Singapore. Association for Computational Linguistics.
- Weiwen Xu, Yang Deng, Wenqiang Lei, Wenlong Zhao, Tat-Seng Chua, and Wai Lam. 2022. [ConReader: Exploring implicit relations in contracts for contract clause extraction](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2581–2594, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.